
Lecture 2

High-Speed I/O

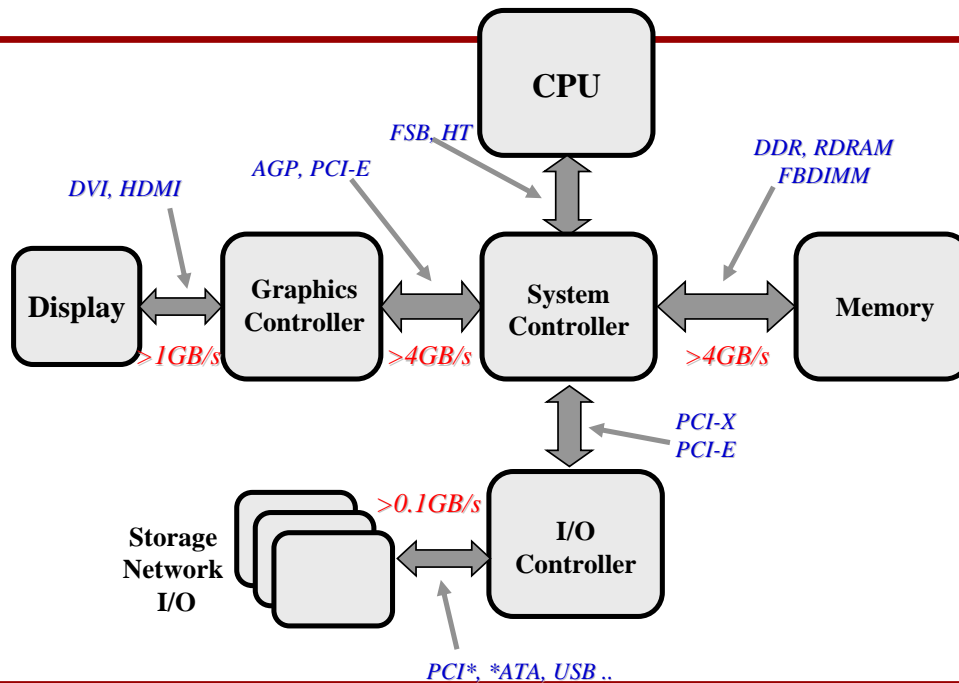
Mark Horowitz
Computer Systems Laboratory
Stanford University
horowitz@stanford.edu

Copyright © 2007 by Mark Horowitz, with material from
Stefanos Sidiropoulos, and Vladimir Stojanovic

Readings

-
- Readings
 - Techniques for High-speed Implementation of Nonlinear Cancellation, Sanjay Kasturia and Jack H. Winters
 - Overview:
 - Your project will be the design of a circuit that processes the input data from a high-speed I/O. This processing is generally done in a mixed signal manner today, but your job will be to build a digital implementation of the algorithm. This lecture will try to give you some background about why I/O rates are important, and what issues need to be resolved to achieve high performance. The next lecture will discuss the operation of the circuit you need to build.

Computers Today



M Horowitz

EE371 Lecture 2

3

Speed of Light: The Difference Between I/O and On-Chip Wires

- First question:
 - Why is I/O different from on-chip wires?
 - Both send signals to each other
- Gates send data to each other all the time
 - Don't generally worry about signals, or delay
 - Model the connection between gates as a capacitor
 - Sometimes a capacitor/resistor network
- Answer:
 - On-chip, ignore the speed of light, assume "c" infinite
 - For external wires can't make that assumption
 - Wire connecting the pins is not an equipotential
 - References are different

M Horowitz

EE371 Lecture 2


4

Finite Speed of Light Ramifications


- Signals must have delay in reaching destination
 - $T_d = L/v$, bits arrive at a different time than when sent
 - Thus must determine 'right' time to sample them
- Wires store energy
 - Current is set by the geometry of wire (what else?)
 - Signal can't see termination resistor (causality)
 - V/I for the line is called the impedance, $Z < 300 \Omega$
 - When signal is traveling on the wire
 - Power goes into the wire before it hits load
 - Since energy is conserved, wire must be storing energy
- Signal is ALWAYS a pair of currents

Link Issues

- Signaling: getting the bit to the receiver






- Timing: Determining which bit is which




Transmission Lines

- Wire where you notice 'c' is finite
 - Current flows in one terminal
 - And flows out the other
- Energy is stored in E and B fields
 - But can model with L, C



Problems : Material Loss




- PCB Loss : skin & dielectric loss
 - Skin Loss $\propto \sqrt{f}$
 - Dielectric loss $\propto f$: a bigger issue at higher f

Dealing With Current Return/References

- Wire Utilization:


- Single Ended
shared signal return path




- Differential
explicit signal return path



- “Pseudo” Differential



Transmission Lines




Two constraints govern behavior at any junction:


- Voltage are equal
 - They are electrically connected
- Power is conserved
 - Energy flow into junction is equal to transmitted and reflected

High-Speed Wires Are Point to Point

- Can't split a wire to go to two location
 - You will get a reflection from the junction
 - Z_1 will see impedance discontinuity



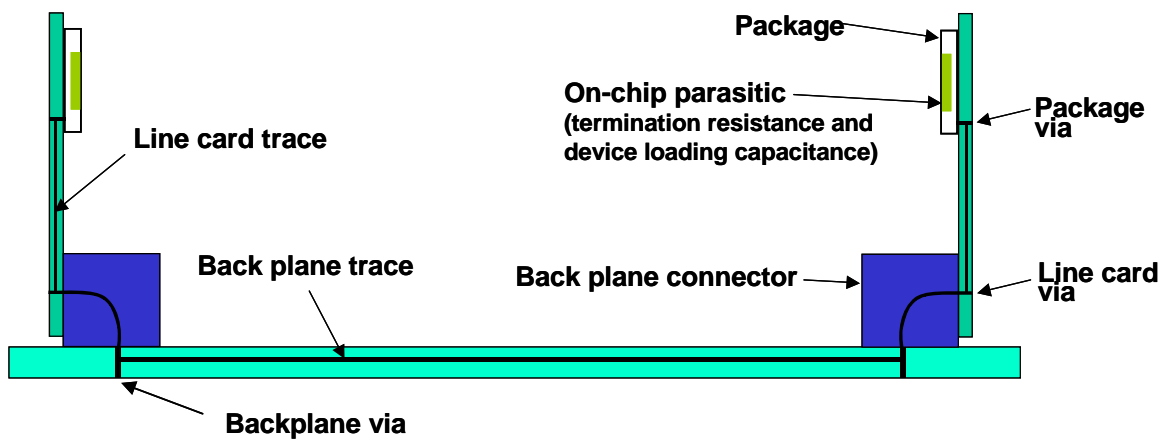
At High Speeds, Vias are Stubs



Top layer signaling
results in large via stub

- Signal energy splits at via
 - If via is short can be modeled as a cap load
 - Causes a reflection in signal
- Higher the frequency, the more sensitive you are to stubs


Backplane Environment



- Line attenuation
- Reflections from stubs (vias)


Backplane Channel

- Loss is variable
 - Same backplane
 - Different lengths
 - Different stubs
 - Top vs. Bot
- Attenuation is large
 - >30dB @ 3GHz
 - But is that bad?




Inter-Symbol Interference (ISI)

- Channel is low pass
 - Our nice short pulse gets spread out



- Dispersion – short latency (skin-effect, dielectric loss)
- Reflections – long latency (impedance mismatches – connectors, via stubs, device parasitics, package)


ISI



- Middle sample is corrupted by
 - 0.2 trailing ISI (from the previous symbol),
 - 0.1 leading ISI (from the next symbol) resulting in 0.3 total ISI
- As a result middle symbol is detected in error



Equalization For Loss :

Goal is to Flatten Response



- Channel is band-limited
- Equalization : boost high-frequencies; or attenuate low freq

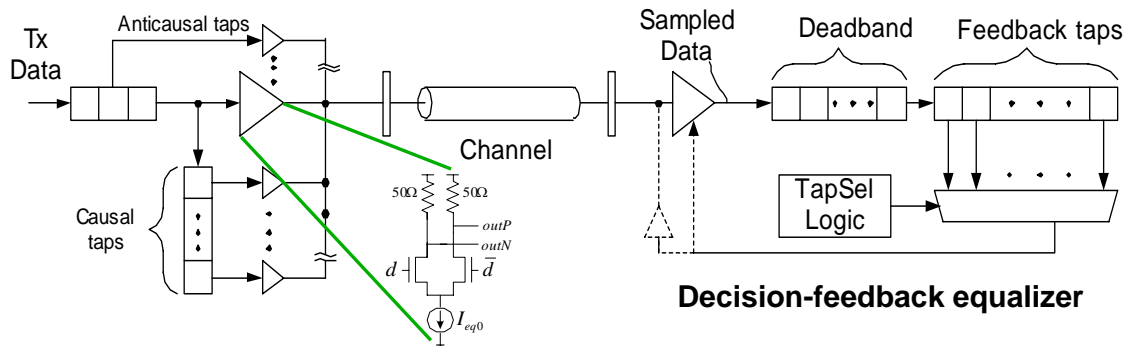
Equalization Mechanisms



- Tx equalization
 - Pre-filter the pulse with the inverse of the channel
 - Filters the low freq. to match attenuation of high freq.
- Rx feedback equalization
 - Subtract the error from the signal

Removing ISI


Linear transmit equalizer



- Transmit and Receive Equalization
 - Changes signal to correct for ISI
 - Initial work was at transmitter

J. Zerbe et al, "Design, Equalization and Clock Recovery for a 2.5-10Gb/s 2-PAM/4-PAM Backplane Transceiver Cell," *IEEE Journal Solid-State Circuits*, Dec. 2003.



Transmit Equalization – Headroom Constraint



- Transmit DAC has limited voltage headroom
- Unknown target signal levels
 - Harder to make adaptive equalization work
- Need to tune the equalizer and receive comparator levels
 - If you have multi-level signals


Removing Interference at Receiver

- Could also build a linear filter
 - Could have gain in the filter
 - But either it would need to be analog and have gain
 - Or need high-speed A/D
 - And real multiplication
 - Sum ($a_i \cdot x_i$)
 - Increases channel noise too




High Frequency Channel Noise: Crosstalk

- Many sources
 - On-chip
 - Package
 - PCB traces
 - Inside connector
- Differential signaling can help
 - Minimize xtalk generation & make effects common-mode
- Both NEXT & FEXT
 - NEXT very destructive if RX and TX pairs are adjacent
 - Full swing-TX coupling into attenuated RX signal
 - Effect on SNR is multiplied by signal loss
 - Simple solution : group RX/TX pairs in connector
 - NEXT typically 3-6%, FEXT typically 1-3%



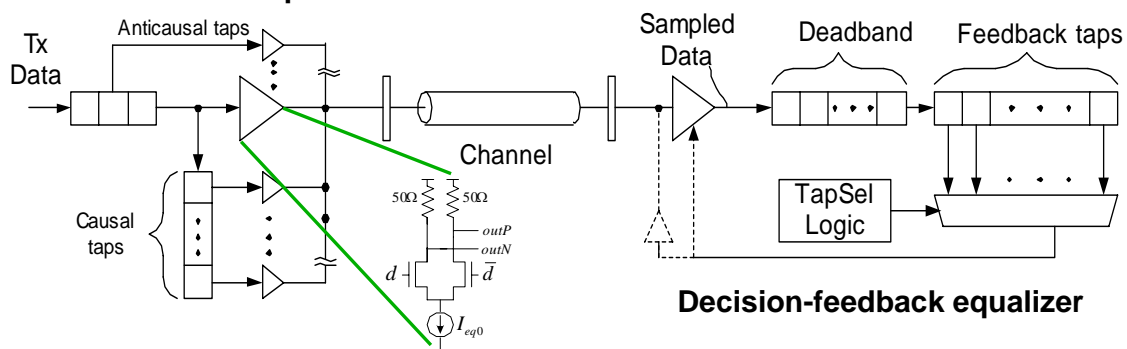
Subtract Out Residual Interference

- Called Decision feedback equalization (DFE)
 - Subtracts error from input
 - No attenuation
- Problem with DFE
 - Need to know interfering bits
 - ISI must be causal
 - Problem - latency in the decision circuit
 - Receive latency + DAC settling < bit time
 - Can increase allowable time by loop unrolling
 - Receive next bit before the previous is resolved



Removing ISI

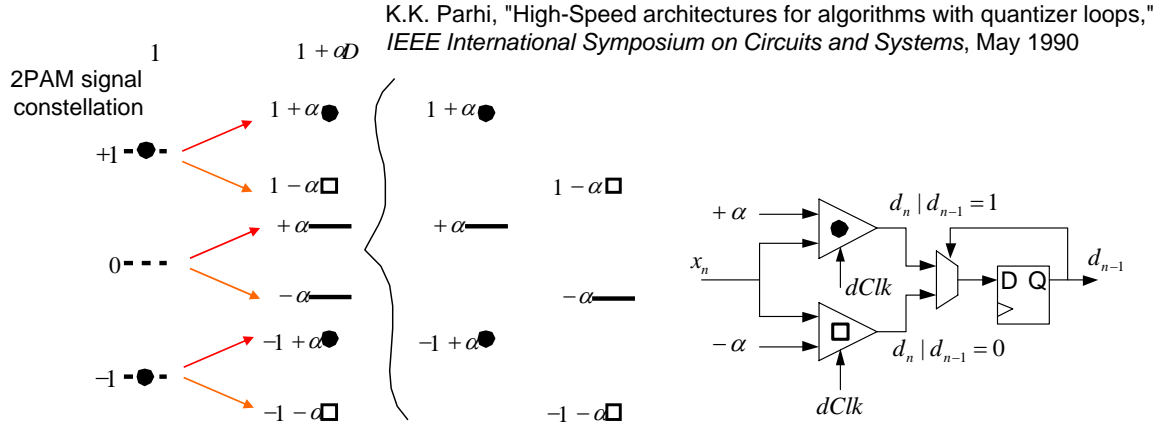
Linear transmit equalizer



- Transmit and Receive Equalization
 - Changes signal to correct for ISI
 - Initial work was at transmitter

J. Zerbe et al, "Design, Equalization and Clock Recovery for a 2.5-10Gb/s 2-PAM/4-PAM Backplane Transceiver Cell," *IEEE Journal Solid-State Circuits*, Dec. 2003.


One Bit Loop Unrolling (for 2 level signal)



- Instead of subtracting the error
 - Move the slicer level to include the interference
 - Slice for each possible level, since previous value unknown

More Bits/Hz

- Multi-level signaling (aka PAM)
 - Convert extra voltage margin to more bits




- Works well when the noise is small
 - Need even more signal processing

Internal Speed Limitation

- Links need good quality clocks with low jitter
 - That means you want them to settle to both Vdd, and Gnd
 - If you make the clock too fast, it will not “rail”
 - And that means it will be prone to jitter
- So one limitation for links is internal clock rate
 - For power efficiency want FO on clock to be around 4
 - Need pulse width 3-4 times the slowest gate
 - Gives around 8 FO4 clock
- For higher speed bit rates
 - Need to generate multiple bits/clock
 - Use non-static CMOS clock circuits (CML & inductors)



Simple Demultiplexing Receiver



- 2-1 demux at the input
- Preconditioning stage: filter/integrate, can be clocked to avoid ISI
 - Reject CM
 - Sometimes not used
- Latch makes decision (4-FO4)

Simple Multiplexing Transmitter

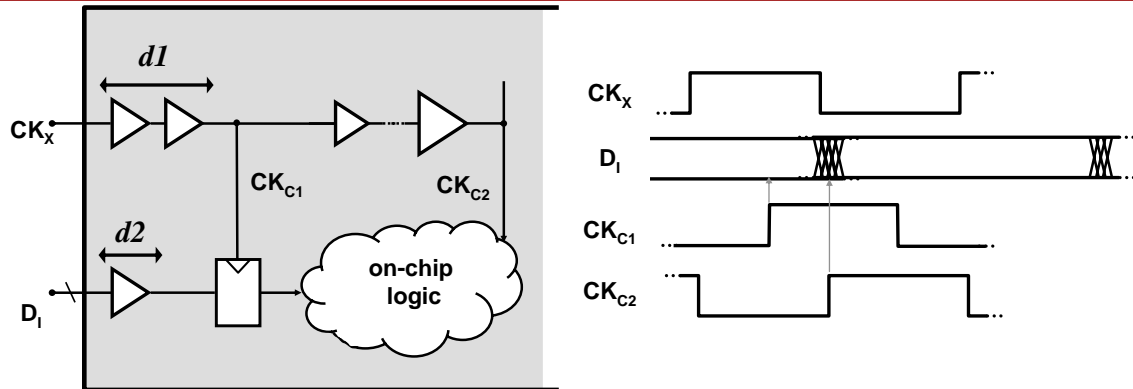
- DDR: send a bit per clock edge
- Critical issues:
 - 50% duty cycle
 - $T_{bit} > 4 \cdot FO4$



I/O Clocking Issues

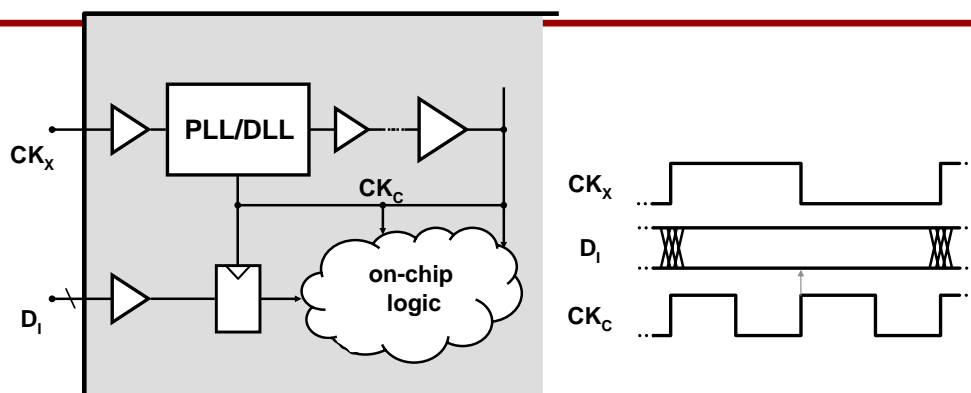
- Remember the clocking issues:
 - Long path constraint (setup time)
 - Short path constraint (hold time)
 - Need to worry about them for I/O as well
- For I/O need to worry about a number of delays
 - Clock skew between chips
 - Data delay between chips
 - Can be larger than a clock cycle (speed of light)
 - Clock skew between external clock and internal clock
 - This can be very large if not compensated
 - It is essentially the insertion delay of the clock tree

System Clocking: Simple Synchronous Systems




- Long bit times compared to on chip delays:
 - Rely on buffer delays to achieve adequate timing margin

PLLs: Creating Zero Delay Buffers




- On-chip clock might be a multiple of system clock:
 - Synthesize on-chip clock frequency
- On-chip buffer delays do not match
 - Cancel clock buffer delay

Used to Argue About PLLs vs DLLs



- Second/third order loop:
 - ➔ Stability is an issue
 - ➔ Frequency synthesis easy
 - ➔ Ref. Clk jitter gets filtered
 - ➔ Phase error accumulates




- First order loop:
 - ➔ Stability ~~guaranteed~~
 - ➔ Frequency synthesis problematic
 - ➔ Ref. Clk jitter propagates
 - ➔ Phase error does not accumulate

After Many Years of Research


- And many papers and products
- One can mess up either a DLL or PLL
 - Each has its own strengths and weaknesses
- If designed correctly, either will work well
 - Jitter will be dominated by other sources
- Many good designs have been published
 - It is now a building block that is often reused
 - We all have our favorites, mine is the dual-loop design
- And yes, people use ring oscillators
 - Still an open question about how much LC helps (in system)

Clocking Structures


- Synchronous:
Same frequency and phase
 - Conventional buses





- Mesochronous
Same frequency, unknown phase
 - Fast memories
 - Internal system interfaces
 - MAC/Packet interfaces



- Plesiochronous:
Almost the same frequency
 - Mostly everything else today

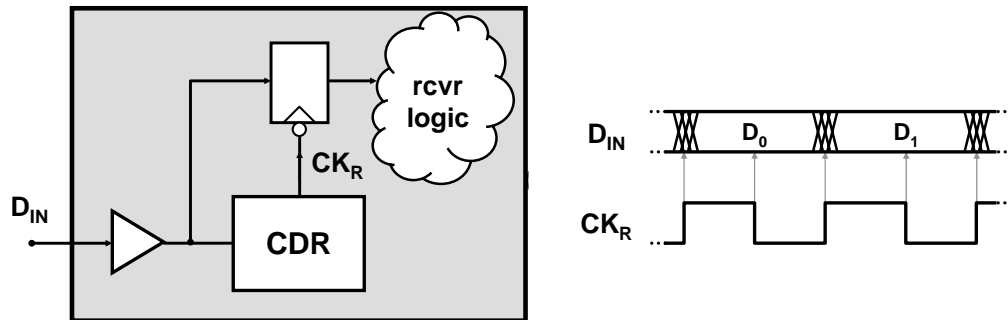


Source Synchronous Systems



- Position on-chip sampling clock at the optimal point
i.e. maximize “timing” margin

Serial Link Circuit



- Recover incoming data fundamental frequency
- Position sampling clock at the “optimal” point